



Multi-PDF RAG Assistant

Upload multiple PDF documents and ask questions across the combined knowledge base.

Upload PDF documents



FinVector-Market.pdf
0.6MB



Machine...g (ML).pdf
476.9KB



sample-1-page.pdf
3.2KB



3 PDF document(s) uploaded.



Uploaded Documents



FinVector-Market.pdf (580,927 bytes)



Machine Learning (ML).pdf (488,316 bytes)



sample-1-page.pdf (3,312 bytes)



Knowledge Base

PDF Documents

3

PDF Pages

19

Indexed Chunks

68

Vector Index

FAISS

Baseline RAG: all uploaded PDFs are combined into one semantic knowledge base.



Indexing Performance



Preview Generated Chunks



Ask Your Documents

Enter your question

tell about distil



Ask

Retrieved 4 candidate chunks.

>  Retrieved Context (Top Candidates)

Answer

Distill is managed by Shan Carter, Chris Olah, and Arvind Satyanarayan. It focuses on making machine learning (ML) and artificial intelligence (AI) more accessible to readers by presenting research in appealing, interactive data visualizations. Distill serves as a neutral platform for multiple authors to publish collaboratively, and its articles are peer-reviewed, appearing in Google Scholar. Additionally, Distill is registered with the Library of Congress and CrossRef.

Sources

 Machine Learning (ML).pdf — Page 2

 FinVector-Market.pdf — Page 3

 FinVector-Market.pdf — Page 10

Query Performance

Retrieval

1.06s

LLM Generation

4.74s

Total Query

5.85s

▼  RAG Architecture

Multiple PDFs



PDF Page Loading



Page Metadata



Recursive Text Splitting



OpenAI Embeddings



FAISS Vector Store



Semantic Retrieval



Retrieved Context



ChatOpenAI



Grounded Answer



PDF + Page Source Citations